


Embeddings and neural text analysis

Basic sequence network template

- Discrete state y_t replaced with continuous vector states h_t ; usually $h_0 = \vec{0}$
- The transition network inputs previous state h_{t-1} and current input x_t (also continuous form) and outputs current state h_t
- y_t for end application derived from h_t using an output network ending in a softmax

Fr



Basic recurrent network (RNN)

- The transition network and the output network are standard multi-layer neural networks with linear combinations and nonlinearities at every hidden layer

Elman network:

$$h_t = \sigma(x_t W_h + h_{t-1} U_h + b_h)$$

tanh *linear combo* *bias vect*

$$y_t = \sigma(h_t W_y + b_y) \quad \text{— logistic regression}$$

softmax

Jordan network:

$$h_t = \sigma(x_t W_h + y_{t-1} U_h + b_h)$$
$$y_t = \sigma(h_t W_y + b_y)$$

- Weights W_h, U_h, b_h of transition networks at all positions are tied; same for weights W_y, b_y in the output network
- RNNs are Turing complete
- Stream of consciousness
- Potential problems: vanishing or exploding gradient
 - I grew up in France ... I speak fluent French

Gated recurrent unit (GRU) [2]

- ▶ Logistic **gates** that decide what fraction of input and previous state to retain/forget

sigmoid

Reset $r_t = \sigma(x_t W_r + h_{t-1} U_r + b_r)$ *Elementwise prod.*

Combine $c_t = \tanh(x_t W_h + (r_t \circ h_{t-1}) U_h + b_h)$

Update $z_t = \sigma(x_t W_z + h_{t-1} U_z + b_z)$

Output $h_t = (1 - z_t) \circ h_{t-1} + z_t \circ c_t$

- ▶ σ is sigmoid; other nonlinearities specified explicitly
- ▶ r_t decides if c_t is computed as in a basic RNN, or h_{t-1} is discarded and only x_t is used
- ▶ z_t decides how to blend c_t and h_{t-1} as-is
- ▶ If we set by force $r_t \approx 1$ and $z_t \approx 1$, equivalent to basic RNN
- ▶ <https://arxiv.org/abs/1412.3555>

Long short-term memory (LSTM) [3]

- ▶ Slightly different arrangement of gates

$$\text{In} \quad i_t = \sigma(x_t U_i + h_{t-1} W_i + b_i)$$

$$\text{Forget} \quad f_t = \sigma(x_t U_f + h_{t-1} W_f + b_f)$$

$$\text{Out} \quad o_t = \sigma(x_t U_o + h_{t-1} W_o + b_o)$$

$$\text{RNN} \quad g_t = \tanh(x_t U_g + h_{t-1} W_g + b_g)$$

$$\text{Memory} \quad c_t = c_{t-1} \circ f_t + g_t \circ i_t$$


$$\text{State} \quad h_t = \tanh(c_t) \circ o_t$$

- ▶ Accuracy comparable to GRUs
- ▶ Widely used for sequence labeling, sequence-to-sequence comparison, sequence-to-sequence translation, etc.
- ▶ Usually left-to-right and right-to-left LSTMs in tandem as a **bi-LSTM**; captures context from both sides
- ▶ Also see [understanding LSTMs, unreasonable effectiveness](#)

How the states h_t are used (c_t usually not)

- ▶ h_t is a compressed representation of the sequence $x_{\leq t}$
- ▶ I.e., the left context of position t
- ▶ Sometimes written \overrightarrow{h}_t
- ▶ To get compressed representation of right context, use a separate right-to-left LSTM
- ▶ Its states are called \overleftarrow{h}_t
- ▶ Together, $\overleftrightarrow{h}_t = [\overrightarrow{h}_t, \overleftarrow{h}_t]$ give a compressed representation of the whole text tailored to position t
- ▶ Consider sequence-to-discrete label classification, e.g., spam, sentiment
- ▶ Feed “average state” $\sum_t \overleftrightarrow{h}_t / T$ as input to a feed-forward network ending in a softmax
- ▶ Given a query, infer if a sentence/passage is relevant
- ▶ Given two sentences, find if one implies the other

Does one sentence imply another?



Quotable expletives and discontent:

“You can't cram the meaning of a whole %&!\$# sentence into a single \$&!#* vector!” — Ray Mooney (2014)

“The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster.”

— Geoffrey Hinton (2014)

RNNs as language models and translators


- ▶ Suppose input is word sequence (x_t)
- ▶ Whose embeddings are inputs to an RNN
- ▶ The label sequence (y_t) can be defined as looking ahead one step in the same word sequence, i.e., (x_{t+1})
- ▶ E.g., having read “I speak fluent ”, what next word do we expect?
- ▶ Thus, an unsupervised corpus can lead to loads of supervised “ y_t ” label data
- ▶ Word embeddings plus LSTM cell model weights constitute “language model”
- ▶ Can hallucinate synthetic text passages
- ▶ Can estimate (comparative) probabilities of given sentences

RNNs as language models and translators (2)

- ▶ Instead of word2vec or GloVe style word embeddings ...
- ▶ ... where 'bank' or 'interest' has only one embedding each
- ▶ Suppose I publish both the LSTM model weights and the word embeddings as trained through the LSTMs
- ▶ Then you run the LSTM through sentences like
 - ▶ He swam close to the river bank
 - ▶ The bank was giving low interest loans
 - ▶ Her interest in music was minimal
- ▶ Then $\overleftarrow{\overrightarrow{h_{\text{bank}}}}$ and $\overleftarrow{\overrightarrow{h_{\text{interest}}}}$ will be quite different in various sentences
- ▶ I.e., context-sensitive word embeddings

RNNs as language models and translators (3)

- ▶ Use of RNNs in translation (sequence to sequence)




- ▶ First encode the English sentence
- ▶ Ending with a special marker '.'
- ▶ "Meaning" of English sentence captured in vector **S**
- ▶ Send in **NULL**, starting the German decoder


Fixed size passage encoding

- ▶ Fixed passage encoding dimension independent of passage length is a problem
- ▶ **Attention** to the rescue

BLEU (gold Fr sent.,
system Fr output)




Encoding the English sentence



$e = (\text{Economic, growth, has, slowed, down, in, recent, years, .})$

Evaluating attention from target position [1]




$e = (\text{Economic, growth, has, slowed, down, in, recent, years, .})$

- ▶ For each target position i
- ▶ Choose one source position j
- ▶ Whose state h_j , along with z_{i-1} informs decoder

Computing attention-weighted average English state

$f = (\text{La, croissance, économique, s'est, ralenti, ces, dernières, années, .})$




$e = (\text{Economic, growth, has, slowed, down, in, recent, years, .})$

- ▶ Choose h_j in a soft manner to support loss backprop
- ▶ Compute a_j as some network function of h_j and z_i
- ▶ Normalize using softmax to add up to 1 and increase skew
- ▶ Compute $\sum_j a_j h_j$ and feed into decoder's $(i + 1)$ th cell

Emitting next French word


$f = (\text{La, croissance, économique, s'est, ralenti, ces, dernières, années, .})$



$e = (\text{Economic, growth, has, slowed, down, in, recent, years, .})$

- ▶ Decoder network itself is not bidirectional
- ▶ Both i and i' can choose a common j^*
- ▶ But each i can (soft-) select only one j

Improved BLEU score



Paper

Attention strength examples

Economic growth has slowed down in recent years .



Das Wirtschaftswachstum hat sich in den letzten Jahren verlangsamt .

Economic growth has slowed down in recent years .



La croissance économique s' est ralenti ces dernières années .

- ▶ *Wirtschafts* means *economic*
- ▶ *Wirtschaftswachstum* means *economic growth*
- ▶ Note the missing line from *growth*
- ▶ Also, *dernières* means *latest*, faint link to *recent*
- ▶ Should allow many-to-many 'fractional' links in general [4]

Self-attention

- ▶ Can we “fake recurrence” by replacing l-to-r or r-to-l recurrence with several layers of (self-)attention?
- ▶ Raw (embedding) inputs x_t
- ▶ Map these to keys k_t , queries q_t , and values v_t
- ▶ Via $k_t \leftarrow W_K x_t + b_K$, $q_t \leftarrow W_Q x_t + b_Q$, $v_t \leftarrow W_V x_t + b_V$
- ▶ For each position t , calculate attention logits over other positions τ via $e_{t,\tau} = q_t \cdot k_\tau$
- ▶ Note that $E = (e_{t,\tau})$ is a $T \times T$ matrix
- ▶ Take softmax as $a_{t,\tau} = \text{softmax}_\tau(e_{t,\tau})$
- ▶ I.e., softmax each row of E : for each t , what is the affinity of the original token at t to each other token?
 - ▶ The **bank** is giving very low **interest rates**
 - ▶ The **drama** proved that **comedy** is **alive** and **well** as a **genre**
- ▶ Compute attention-weighted values as outputs $o_t = \sum_\tau a_{t,\tau} v_\tau$

Self-attention (2)

- ▶ One's output $o_t^{(\ell)}$ is another's input $x_t^{(\ell+1)}$ (rinse and repeat)
- ▶ How to train?
- ▶ Hide ('mask') some fraction of tokens and predict them from o_t
- ▶ (word2vec was doing something similar, predicting one token at a time)
- ▶ Add a gazillion bells and whistles to get ... Transformers!
- ▶ What do BERT, RoBERTa, ALBERT, SpanBERT, DistilBERT, SesameBERT, SemBERT, SciBERT, BioBERT, MobileBERT, TinyBERT and CamemBERT all have in common?

General notation for GRUs and LSTMs

- ▶ In general we can use ‘RNN’ as an abstract cell:

$$\mathbf{h}_t = \text{RNN}(\mathbf{h}_{t-1}, x_{t-1})$$

- ▶ \mathbf{h}_t is the recurrent state and x_t is the input
- ▶ We regard the input x_t as 1-of- Σ discrete, written as a 1-hot vector, and looked up in an embedding matrix to send a vector into the RNN cell
- ▶ A separate output network converts states to distributions over the output¹ vocabulary Σ , from which output symbol y_t can be sampled

$$y_t \sim \text{OUT}(\mathbf{h}_t) \in \Delta^\Sigma,$$

where Δ^Σ is the multinomial distribution over $|\Sigma|$ outcomes

- ▶ Standard RNNs and GRUs fit this exactly
- ▶ LSTMs pass *two* vectors $\mathbf{c}_t, \mathbf{h}_t$ from each position to next
- ▶ These can be notationally clubbed into a single \mathbf{h}_t

General notation for GRUs and LSTMs (2)

- ▶ And OUT ignores the c_t part
- ▶ Or, we can explicitly write

$$[\mathbf{c}_t, \mathbf{h}_t] = \text{LSTM}([\mathbf{c}_{t-1}, \mathbf{h}_{t-1}], x_{t-1})$$
$$y_t \sim \text{OUT}(\mathbf{h}_t)$$

- ▶ In language modeling applications we have ‘gold’ y_t^* available, which is just the next word x_t
- ▶ Our goal there is to maximize $\text{OUT}(\mathbf{h}_t)[x_t]$, or the overall log likelihood

$$\sum_t \log \text{OUT}(\mathbf{h}_t)[x_t]$$

- ▶ In a seq2seq translation network without attention, there is an encoder RNNe and a decoder RNNd
- ▶ During encoding there is no output
- ▶ Let the input sequence have length N and output sequence have length M

General notation for GRUs and LSTMs (3)

- ▶ The last encoder state is \mathbf{h}_N
- ▶ To avoid confusion name the states of the decoder \mathbf{g}_m ; $\mathbf{g}_0 \equiv \mathbf{h}_N$
- ▶ And let the inputs and outputs of the decoder be y_m
- ▶ For the decoder, $y_0 = \text{START}$ is a special marker token
- ▶ The decoder does have an associated OUT network

$$\mathbf{g}_m = \text{RNND}(\mathbf{g}_{m-1}, y_{m-1})$$

$$y_m \sim \text{OUT}(\mathbf{g}_m)$$

- ▶ Here we have a gold translated sequence $y_m^*, m = 1, \dots, M$
- ▶ 0/1 loss is like $\sum_m \llbracket y_m \neq y_m^* \rrbracket$
- ▶ Is $\sum_m -\log \text{OUT}(\mathbf{g}_m)[y_m^*]$ an ok surrogate?
- ▶ In the basic [attention model](#), every output position m is 'explained'/supported by an input position n , chosen in a soft manner

General notation for GRUs and LSTMs (4)

- ▶ Encoder RNN_e works as before, but all $\mathbf{h}_1, \dots, \mathbf{h}_N$ can be accessed by decoder RNN_d
 - ▶ Assume RNN_d is a GRU decoder for simplicity
 - ▶ Attention from output position m on input position n may be modeled as $a_{mn} \propto \exp(\mathbf{g}_m \cdot \mathbf{h}_n)$
 - ▶ Normalize to multinomial distribution for each m :
$$a_{mn} = \exp(\mathbf{g}_m \cdot \mathbf{h}_n) / \sum_{n'} \exp(\mathbf{g}_m \cdot \mathbf{h}_{n'})$$
 - ▶ Attention-weighted source state selection for output position m will be $\hat{\mathbf{h}}_m = \sum_n a_{mn} \mathbf{h}_n$
 - ▶ Now decoder recurrence is set up as
$$\mathbf{g}_m = \text{RNNd}(\mathbf{g}_{m-1}, \hat{\mathbf{h}}_m, y_{m-1})$$
 - ▶ And output $y_m \sim \text{OUT}(\mathbf{g}_m)$ is sampled as before
 - ▶ (Mild variations used in the literature)
- ▶ HW Practice writing vectorized Tf/Pytorch code for above

¹Input and output vocabularies may be same or different

Summary of using neural networks for text

- ▶ Start with individual word embeddings
- ▶ ConvNets provide somewhat position-sensitive aggregation of words of a sentence or passage, without demanding hard parses
- ▶ If parses are available, can compose word embeddings along the parse structure
- ▶ (Not all compositions are valid, e.g., “New York”, “rat race”, “red carpet”, “snake oil”)
- ▶ Counterpart to chain CRFs is the RNN family
- ▶ Exploding and vanishing gradients
- ▶ Addressed by GRU and LSTM
- ▶ Applied in parsing, entailment testing, machine translation, question answering, ...
- ▶ Various attention mechanisms boost performance